

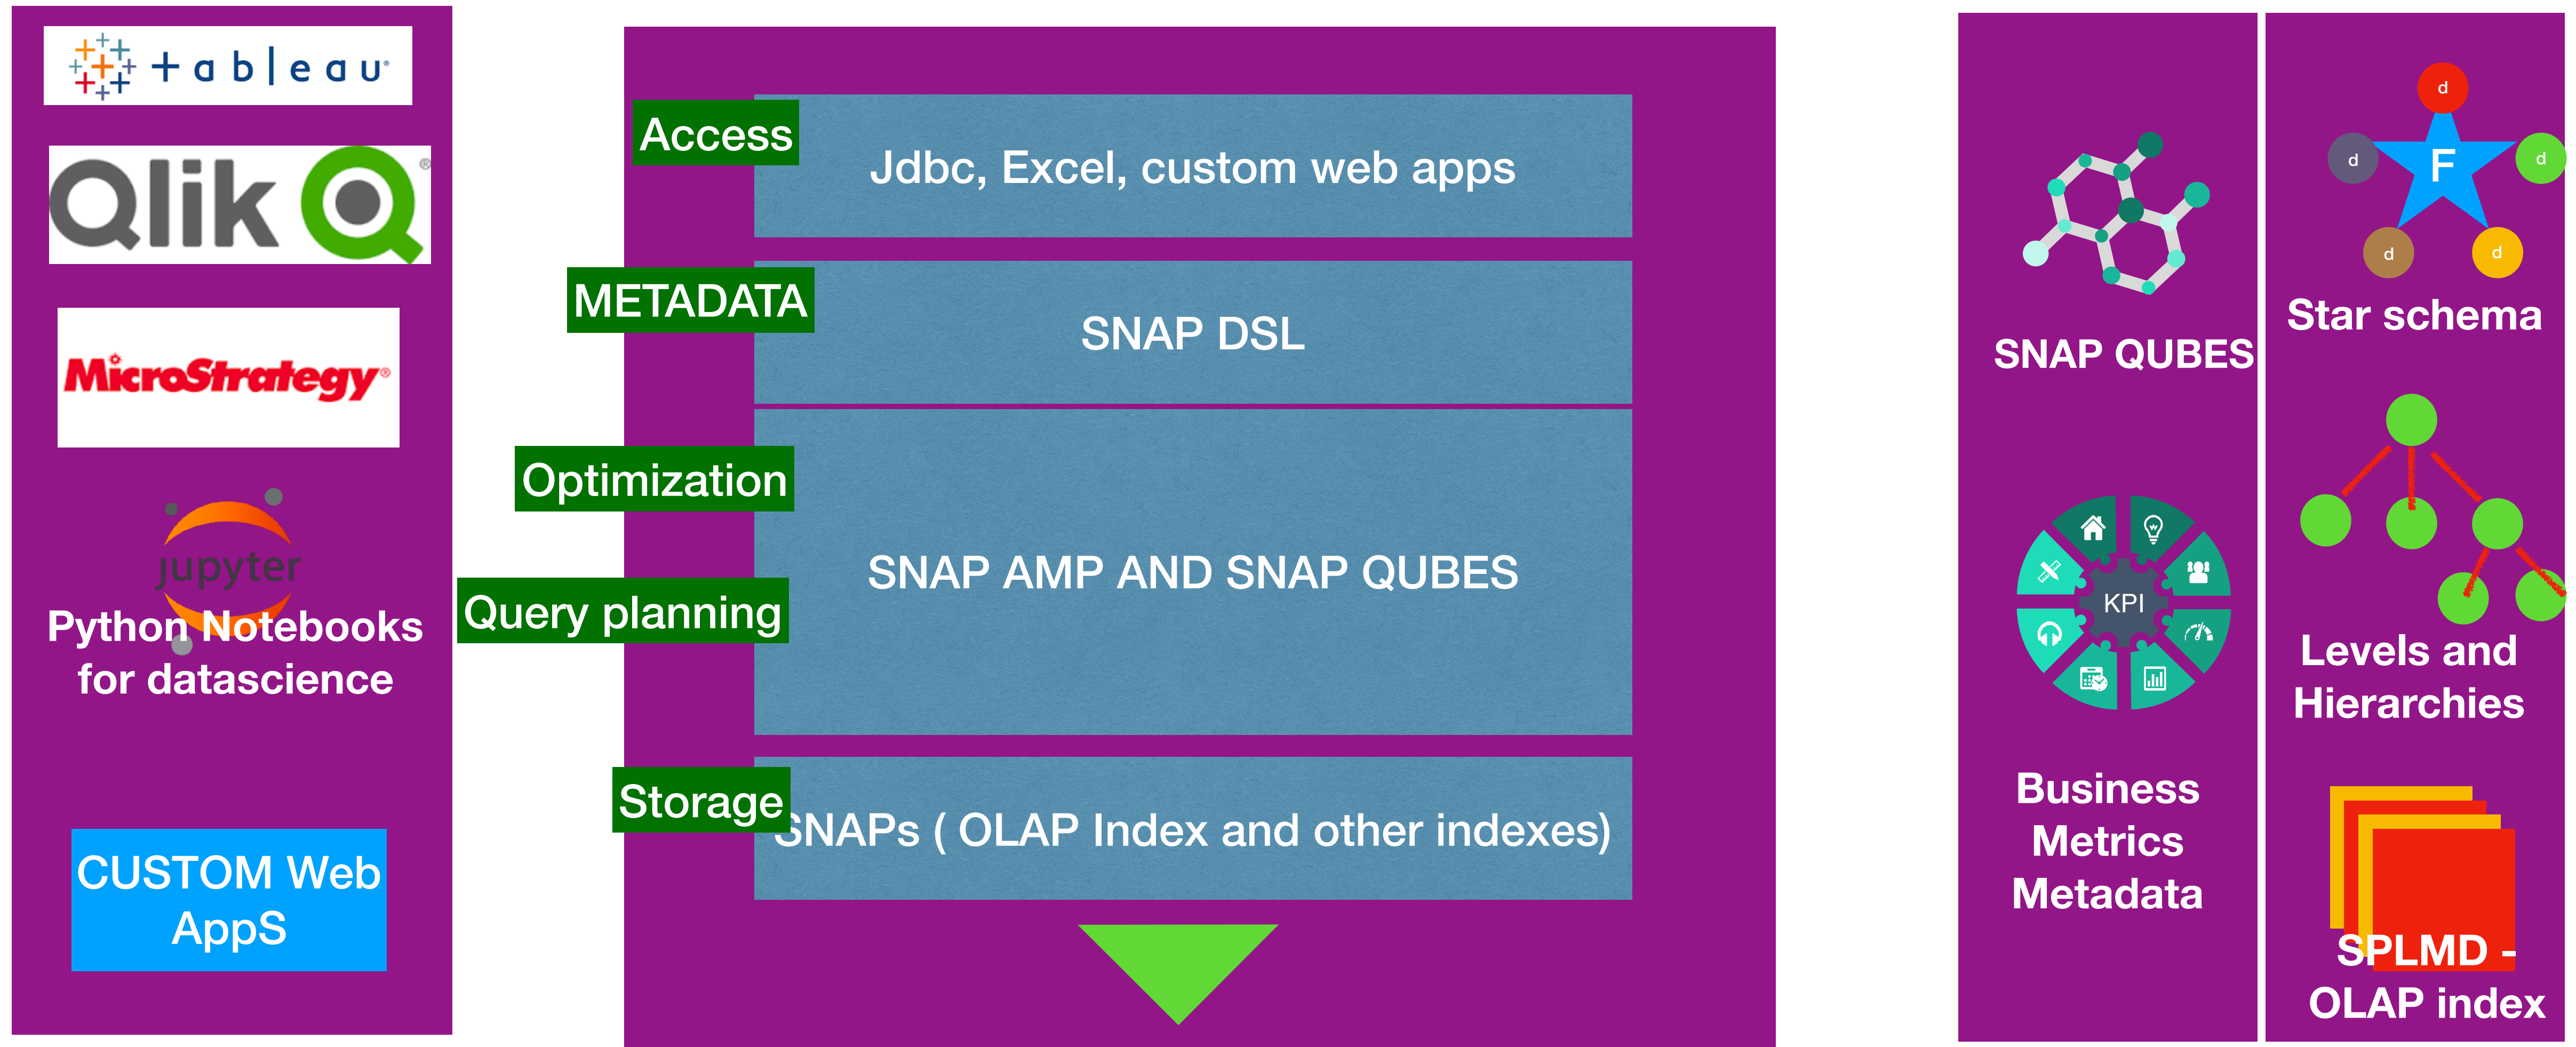
Sparkline SNAP

Enterprise Datamart | Sub second queries

JUNE, 2017 |



Solution Architecture



Set up an index

SNAP indexes

create external table

create olap index

spark / hive external tables

Logical Schema

SNAP OLAP Index

compressed columnar data and index

```
0 Jan 27 15:33 SUCCESS
884 Jan sales_demo p_year=1992
884 Jan 29 08:10 p_year=1993
884 Jan 29 08:17 p_year=1994
884 Jan 29 08:23 p_year=1995
884 Jan 29 08:29 p_year=1996
884 Jan 29 08:36 p_year=1997
612 Jan 29 08:40 p_year=1998
```

physical data in HDFS /S3

```
0 Jan 27 15:33 SUCCESS
884 Jan sales_snap p_year=1992
884 Jan 29 08:10 p_year=1993
884 Jan 29 08:17 p_year=1994
884 Jan 29 08:23 p_year=1995
884 Jan 29 08:29 p_year=1996
884 Jan 29 08:36 p_year=1997
612 Jan 29 08:40 p_year=1998
```



SNAP Qubes - Step 1

Set up Spark datasource table


```
CREATE TABLE IF NOT EXISTS sales_demo
(
  o_orderkey      INTEGER,
  o_orderstatus   STRING,
  o_totalprice     DOUBLE,
  o_orderdate      STRING,
  o_orderpriority STRING,
  o_shippriority  INTEGER,
  l_linenumbers   INTEGER,
  l_quantity      DOUBLE,
  l_extendedprice DOUBLE,
  l_discount       DOUBLE,
  l_tax            DOUBLE,
  l_returnflag     STRING,
  l_linestatus     STRING,
  l_shipdate       STRING,
  l_commitdate     STRING,
  l_receiptdate    STRING,
  l_shipmode       STRING,
  order_year       STRING,
  ps_availqty      INTEGER,
  ps_supplycost    DOUBLE,
  s_name           STRING,
  s_acctbal         DOUBLE,
  s_nation          STRING,
  s_region          STRING,
  p_name           STRING,
  p_mfgr           STRING,
  p_brand           STRING,
  p_type           STRING,
  p_size           INTEGER,
  p_container       STRING,
  p_retailprice     DOUBLE,
  c_name           STRING,
  c_phone           STRING,
  c_acctbal         DOUBLE,
  c_mktsegment      STRING,
  c_nation          STRING,
  c_region          STRING,
  p_year           STRING,
  p_month          STRING
)
```

```
using csv
OPTIONS (path "s3://SNAP/samples/sales_demo_par")
partitioned by ( p_year, p_month )
```

- A. Your Data is in HDFS /S3 or S3 compatible object storage or an external datawarehouse.
- B. Connect to SNAP (using jdbc:hive://<ip address where snap server is running: port and any authentication you may have configured. (see hive URL format for example of authentication)
- C. create an external table pointing at the data
 - A. See example to the left.
- D. if your data is not partitioned, partition the data.

A simple Qube - step 2

o_orderkey	STRING,
o_orderstatus	STRING,
o_orderpriority	STRING,
o_shippriority	STRING,
l_linenum	STRING,
l_returnflag	STRING,
l_linestatus	STRING,
l_shipmode	STRING,
order_year	STRING,
s_name	STRING,
s_nation	STRING,
s_region	STRING,
p_name	STRING,
p_mfgr	STRING,
p_brand	STRING,
p_type	STRING,
p_size	STRING,
p_container	STRING,
c_name	STRING,
c_phone	STRING,
c_acctbal	DOUBLE,
c_mktsegment	STRING,
c_nation	STRING,
c_region	STRING,
o_orderdate	STRING,
l_shipdate	STRING,
l_commitdate	STRING,
l_receiptdate	STRING,
o_totalprice	DOUBLE,
l_quantity	DOUBLE,
l_extendedprice	DOUBLE,
l_discount	DOUBLE,
l_tax	DOUBLE,
s_acctbal	DOUBLE,
ps_availqty	INTEGER,
ps_supplycost	DOUBLE,
p_retailprice	DOUBLE,



**Define columns
As either metrics
Or dimensions**

Dimensions

```
create olap index sales_snap on sales_demo
dimension p_name is not nullable
timestamp dimension l_shipdate spark
timestamp dimension o_orderdate
timestamp dimension l_commitdate
is nullable nullvalue
"1992-01-01T00:00:00.000"
timestamp dimension l_receiptdate
is not nullable
```

dimensions

```
"o_orderkey,o_orderstatus,o_orderpriority,o_
shippriority,l_linenum,l_returnflag,l_lin
estatus,l_shipmode,s_name,s_nation,s_region,
p_mfgr,p_brand,p_type,p_size,p_container,c_n
ame,c_phone,c_mktsegment,c_nation,c_region"
```

metrics

```
"o_totalprice,l_quantity,l_extendedprice,l_d
iscount,l_tax,ps_availqty,ps_supplycost,s_ac
ctbal,p_retailprice,c_acctbal"
OPTIONS ( path "/SNAP/samples/
sales_demo_index", rowFlushBoundary "10000",
rowsPerSegment "50000",
nonAggregateQueryHandling
"push_project_and_filters")
partition by order_year
```

Metrics



**Source table we
defined in the previous step**

©SparklineData Confidential 2018

Qube sections

```
create olap index sales_snap on sales_demo
dimension p_name is not nullable
timestamp dimension l_shipdate spark is nullable nullvalue
"1992-01-01T00:00:00.000"
timestamp dimension o_orderdate
timestamp dimension l_commitdate
is nullable nullvalue "1992-01-01T00:00:00.000"
dimension o_orderstatus is nullable nullvalue "NA"
dimensions
"o_orderkey,o_orderpriority,o_shippriority,l_linenummer,l_
returnflag,..."
metrics "o_totalprice,l_quantity,l_extendedprice"

OPTIONS ( path "/SNAP/samples/sales_demo_index",,
avgSizePerPartition "100mb",
preferredSegmentSize "100mb",
rowFlushBoundary "100000")
partition by order_year
```

Name of the
index

Name of
external hive
or spark table

define rules on
dimensions if any
- Default
nullValue

Define dimensions and metrics(exclude
those you have already defined earlier as
timestamps, in this section

Define index options - see docs for detailed
explanation of options.




Partitioning indexes

```
create olap index sales_snap on sales_demo
dimension p_name is not nullable
timestamp dimension l_shipdate is nullable nullvalue "1992-01-01T00:00:00.000"
timestamp dimension o_orderdate
timestamp dimension l_commitdate
            is nullable nullvalue "1992-01-01T00:00:00.000"

dimensions
"o_orderkey,o_orderstatus,o_orderpriority,o_shippriority,l_linenum,l_returnflag,..."
metrics "o_totalprice,l_quantity,l_extendedprice,..."

OPTIONS ( path "/SNAP/samples/sales_demo_index", rowFlushBoundary "10000",
rowsPerSegment "50000",
    avgSizePerPartition "100mb",
        avgNumRowsPerPartition "1660000",
        preferredSegmentSize "100mb",
        rowFlushBoundary "100000")

partition by order_year
```



Indexes can be partitioned even if the source table is not

If your source data is already partitioned use the same partition scheme for the indexes as well.

when new data is added to your source partition insert data into the index as well

Source table is a flat table

Index is partitioned on "order_year"



adding data to snap indexes

Insert olap index *sales_snap* of *sales_demo*

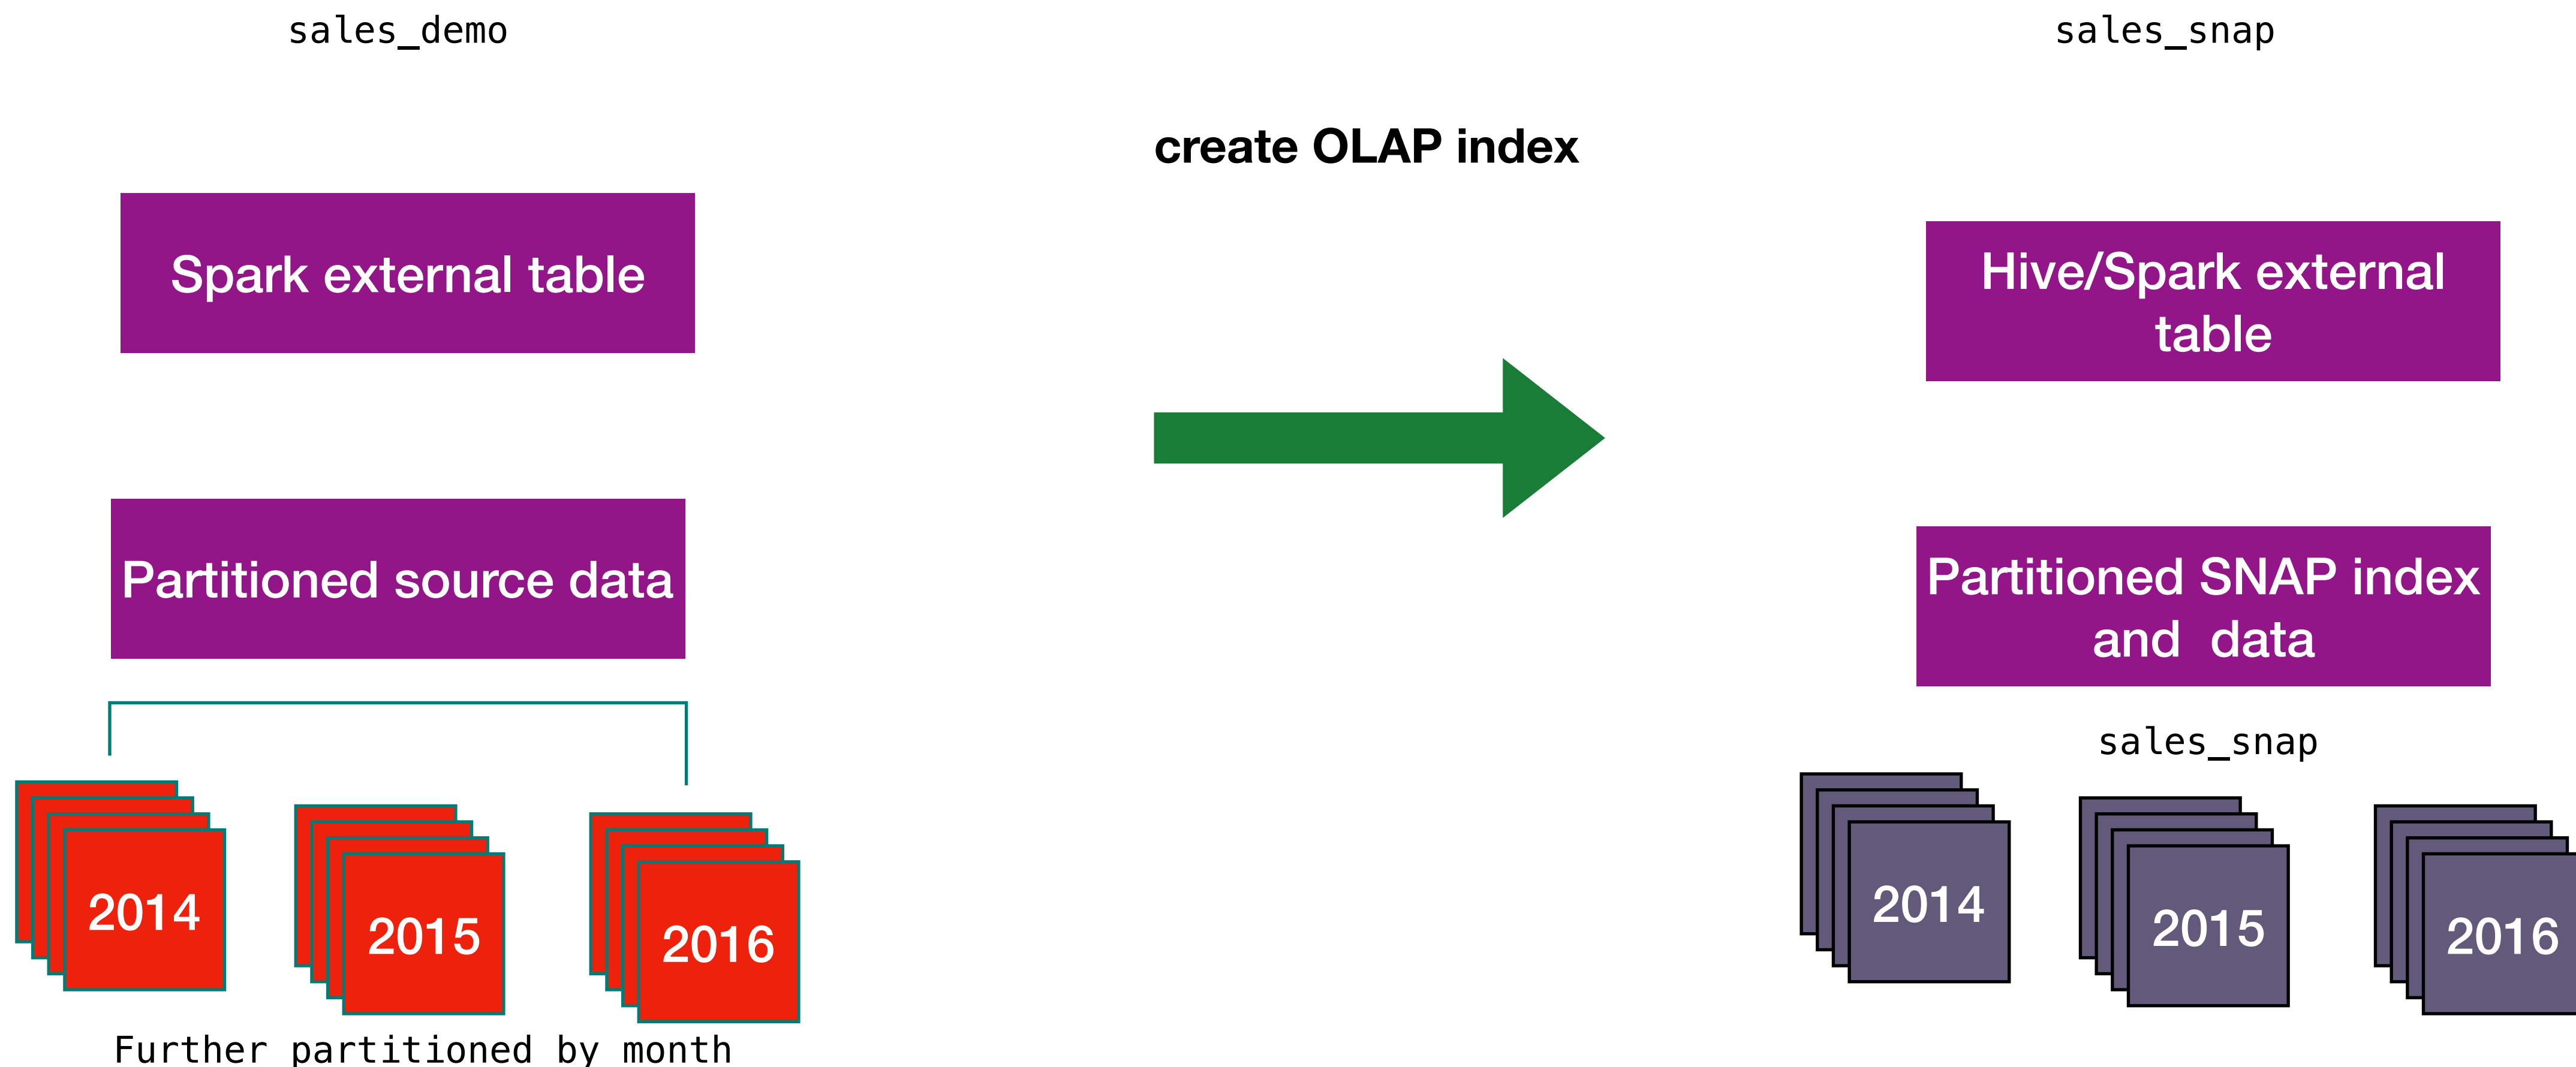
partitions *order_year = "2017"*

Insert **overwrite** olap index *sales_snap* of *sales_demo*

partitions *order_year = "2016"*



Putting it all together



Putting it all together

Insert olap index *sales_snap* of *sales_demo*

sales_demo

partitions *order_year = "2017" order_month = "01"*


sales_snap

Hive/Spark external
table


Hive/Spark external
table

Partitioned source data

Partitioned SNAP index
and data



Further partitioned by month



sales_snap



Putting it all together

Insert overwrite olap index *sales_snap* of *sales_demo*

sales_demo

partitions order_year = "2016" order_month = "03"

sales_snap

Spark external table


Hive/Spark external
table



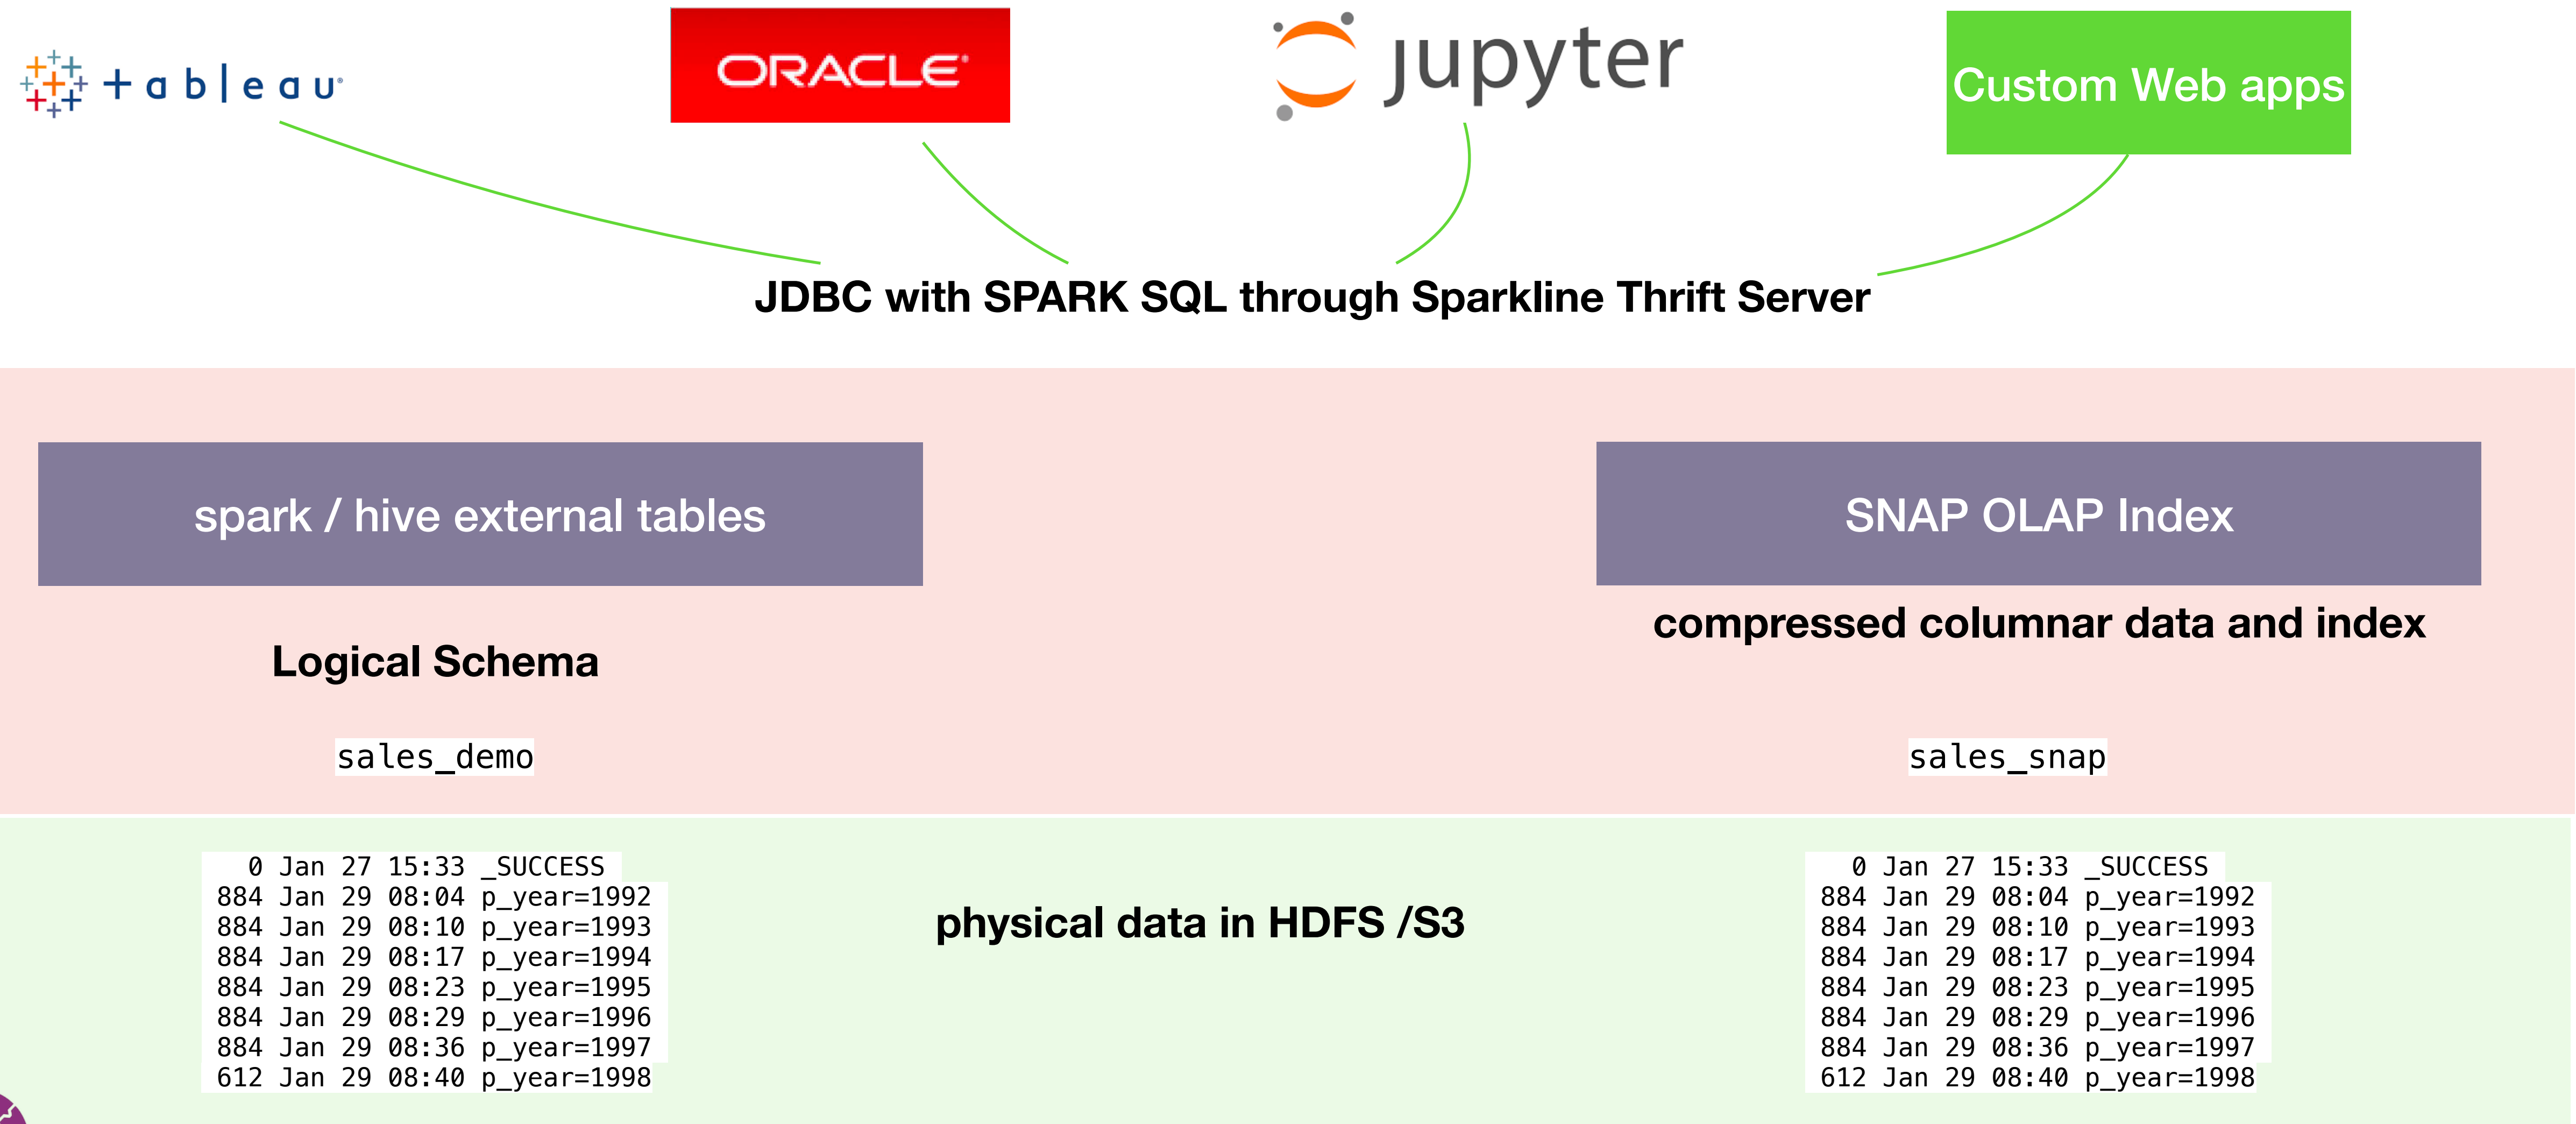
Partitioned source data

Updating data that has changed for previous
Periods

Partitioned SNAP index
and data



SNAP QUERIES



SNAP QUERIES



JDBC with SPARK SQL through Sparkline Thrift Server

select store, sum(sales) from

sales_demo

group by store

spark / hive external tables

Query gets rewritten to
use the SNAP index
Up to 100x faster
for slice and dice queries

SNAP OLAP Index

compressed columnar data and index

sales_snap

physical source data can be archived to a
deep storage like s3
and not needed for queries

```
0 Jan 27 15:33 _SUCCESS
884 Jan 29 08:04 p_year=1992
884 Jan 29 08:10 p_year=1993
884 Jan 29 08:17 p_year=1994
884 Jan 29 08:23 p_year=1995
884 Jan 29 08:29 p_year=1996
884 Jan 29 08:36 p_year=1997
612 Jan 29 08:40 p_year=1998
```

```
0 Jan 27 15:33 _SUCCESS
884 Jan 29 08:04 p_year=1992
884 Jan 29 08:10 p_year=1993
884 Jan 29 08:17 p_year=1994
884 Jan 29 08:23 p_year=1995
884 Jan 29 08:29 p_year=1996
884 Jan 29 08:36 p_year=1997
612 Jan 29 08:40 p_year=1998
```

physical data and index in HDFS /S3



SNAP with Tableau

The screenshot shows the Tableau interface with a 'Spark SQL' connection dialog box open. The dialog box contains the following fields and options:

- Server:** localhost
- Port:** 10000
- Enter information to sign in to the server:**
 - Type:** SparkThriftServer (Spark 1.1 an...)
 - Authentication:** Username
 - Username:** hadoop
- Buttons:** Initial SQL..., Sign In

The background shows the Tableau workspace with a data source named 'sparkline_sales_cube (default.sparkline_sales_cube) (default)' and a table named 'sales_demo_par'. The table has columns: #, sales_demo_par, O Shippriority, #, sales_demo_par, L Linenumbr, #, sales_demo_par, L Quantity.

Start tableau and connect to the spark-sql connector

Choose the host where you started the thrift server and the corresponding port

types of queries in tableau

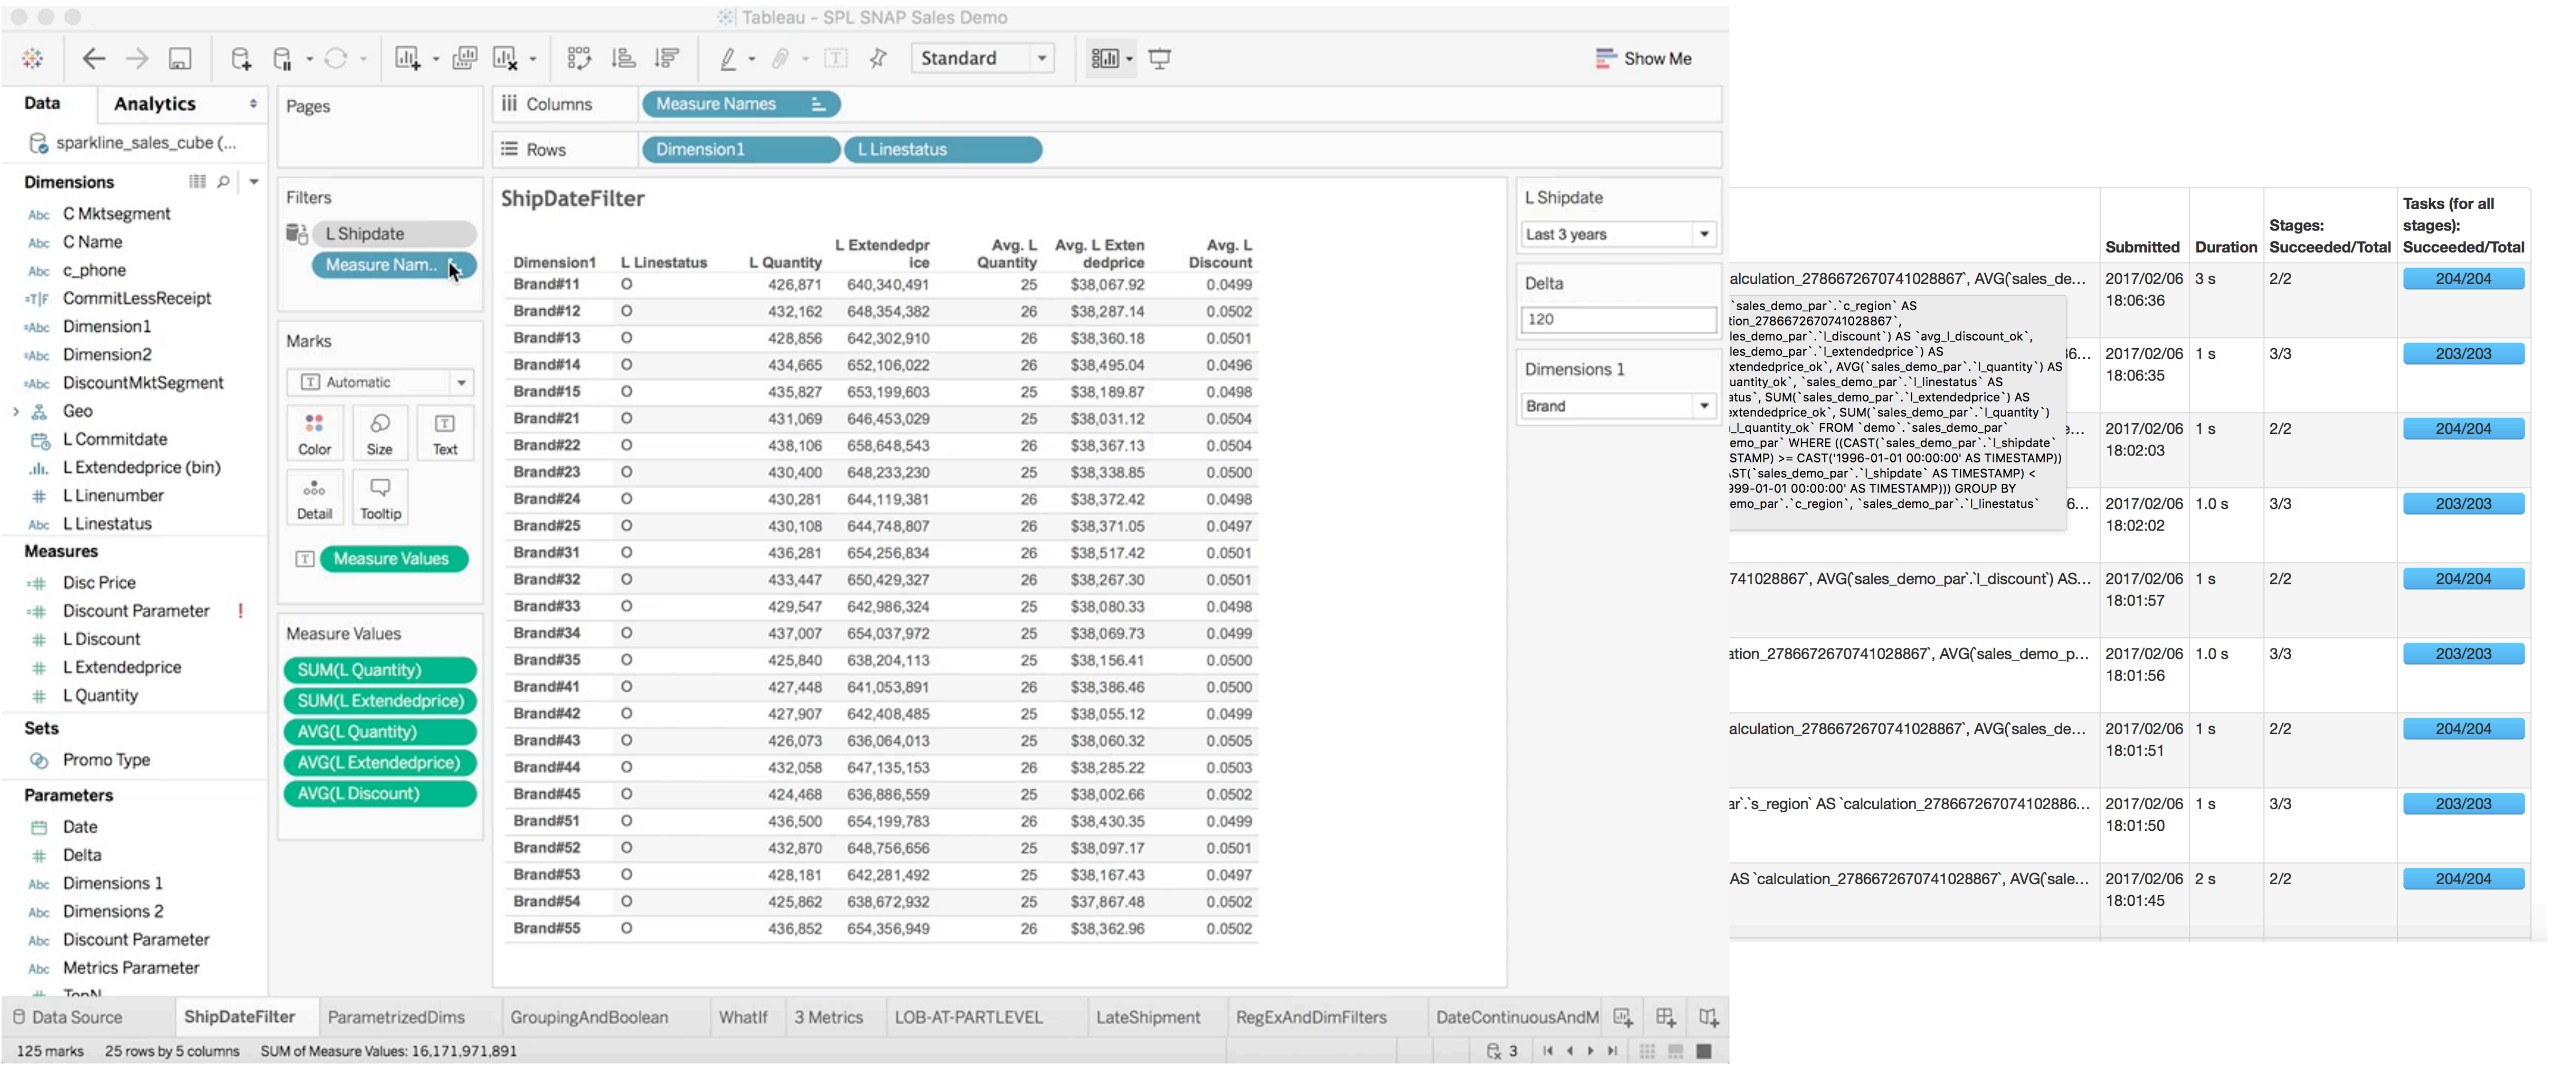
The screenshot shows the Tableau interface for a 'SPL SNAP Sales Demo'. The main view is a table titled 'ShipDateFilter' with columns: Dimension1, L Linestatus, L Quantity, L Extendedprice, Avg. L Quantity, Avg. L Extendedprice, and Avg. L Discount. The table lists categories: AUTOMOBILE, BUILDING, FURNITURE, HOUSEHOLD, and MACHINERY. Annotations highlight two types of filters:

- Timestamp filter:** Points to the 'L Shipdate' field in the Filters shelf and the 'L Shipdate' dropdown menu on the right, which is set to 'Last 3 years'.
- Dimension filter:** Points to the 'Market Segment' dropdown menu in the 'Dimensions 1' section on the right, with the text 'choose drop down to dynamically change the'.

The left sidebar shows the Data pane with Dimensions (C Mktsegment, C Name, c_phone, CommitLessReceipt, Dimension1, Dimension2, DiscountMktSegment, Geo, L Commitdate, L Extendedprice (bin), L Linenumber, L Linestatus) and Measures (Disc Price, Discount Parameter, L Discount, L Extendedprice, L Quantity). The bottom status bar indicates '25 marks', '5 rows by 5 columns', and 'SUM of Measure Values: 16,171,206,864'.

Dimension1	L Linestatus	L Quantity	L Extendedprice	Avg. L Quantity	Avg. L Extendedprice	Avg. L Discount
AUTOMOBILE	O	2,149,526	3,223,659,726	25	\$38,238.51	0.0501
BUILDING	O	2,137,271	3,196,897,192	25	\$38,061.02	0.0501
FURNITURE	O	2,164,445	3,248,668,562	26	\$38,315.64	0.0499
HOUSEHOLD	O	2,142,257	3,211,622,498	25	\$38,200.40	0.0500
MACHINERY	O	2,185,187	3,279,388,938	26	\$38,320.02	0.0501

Example



	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
alculation_2786672670741028867', AVG(sales_de...	2017/02/06 18:06:36	3 s	2/2	204/204
`sales_demo_par`.`c_region` AS tion_2786672670741028867', les_demo_par`.`l_discount`) AS `avg_l_discount_ok`, les_demo_par`.`l_extendedprice`) AS xtendedprice_ok', AVG(`sales_demo_par`.`l_quantity`) AS quantity_ok', `sales_demo_par`.`l_linestatus` AS atus', SUM(`sales_demo_par`.`l_extendedprice`) AS xtendedprice_ok', SUM(`sales_demo_par`.`l_quantity`) l_quantity_ok` FROM `demo`.`sales_demo_par` emo_par` WHERE ((CAST(`sales_demo_par`.`l_shipdate` STAMP) >= CAST('1996-01-01 00:00:00' AS TIMESTAMP)) (ST(`sales_demo_par`.`l_shipdate` AS TIMESTAMP) < 999-01-01 00:00:00' AS TIMESTAMP))) GROUP BY emo_par`.`c_region`, `sales_demo_par`.`l_linestatus`	2017/02/06 18:06:35	1 s	3/3	203/203
	2017/02/06 18:02:03	1 s	2/2	204/204
	2017/02/06 18:02:02	1.0 s	3/3	203/203
741028867', AVG(sales_demo_par`.`l_discount`) AS...	2017/02/06 18:01:57	1 s	2/2	204/204
ation_2786672670741028867', AVG(sales_demo_p...	2017/02/06 18:01:56	1.0 s	3/3	203/203
alculation_2786672670741028867', AVG(sales_de...	2017/02/06 18:01:51	1 s	2/2	204/204
ar`.`s_region` AS `calculation_278667267074102886...	2017/02/06 18:01:50	1 s	3/3	203/203
AS `calculation_2786672670741028867', AVG(sale...	2017/02/06 18:01:45	2 s	2/2	204/204



Behind the scenes

▶ Event Timeline

Completed Jobs (281)

Job Id (Job Group)	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
280 (a2f66739-e7cc-4bf0-9b64-50ee40de0732)	SELECT `sales_demo_par`.`c_region` AS `calculation_2786672670741028867`, AVG(sales_de... run at AccessController.java:-2	2017/02/06 18:06:36	3 s	2/2	204/204
279 (c81e01d6-a2ce-4eb9-9a76-e7139fada9e7)	SELECT * FROM (SELECT `sales_... run at AccessController.java:-2	2017/02/06 18:06:35	1 s	3/3	203/203
278 (eb380f87-1d10-4489-8395-93946f339929)	SELECT `sales_demo_par`.`p_br... run at AccessController.java:-2	2017/02/06 18:02:03	1 s	2/2	204/204
277 (a481623d-f338-4368-93e4-68ca7789927b)	SELECT * FROM (SELECT `sales_... run at AccessController.java:-2	2017/02/06 18:02:02	1.0 s	3/3	203/203
276 (60bf3a02-6074-4421-8faf-7b1a6614b325)	SELECT 'NA' AS `calculation_2786672670741028867`, AVG(sales_demo_par`.`l_discount`) AS... run at AccessController.java:-2	2017/02/06 18:01:57	1 s	2/2	204/204
275 (104842e1-b165-4b0e-b117-66877ed1d7fb)	SELECT * FROM (SELECT 'NA' AS `calculation_2786672670741028867`, AVG(sales_demo_p... run at AccessController.java:-2	2017/02/06 18:01:56	1.0 s	3/3	203/203
274 (b868108c-8c92-41d4-8604-a5f810e318e0)	SELECT `sales_demo_par`.`s_region` AS `calculation_2786672670741028867`, AVG(sales_de... run at AccessController.java:-2	2017/02/06 18:01:51	1 s	2/2	204/204
273 (b89ef74c-3978-4536-96a5-6b6937d1f1e5)	SELECT * FROM (SELECT `sales_demo_par`.`s_region` AS `calculation_278667267074102886... run at AccessController.java:-2	2017/02/06 18:01:50	1 s	3/3	203/203
272 (162af573-c08f-443a-9ba5-e08139acff24)	SELECT `sales_demo_par`.`o_orderstatus` AS `calculation_2786672670741028867`, AVG(sale... run at AccessController.java:-2	2017/02/06 18:01:45	2 s	2/2	204/204

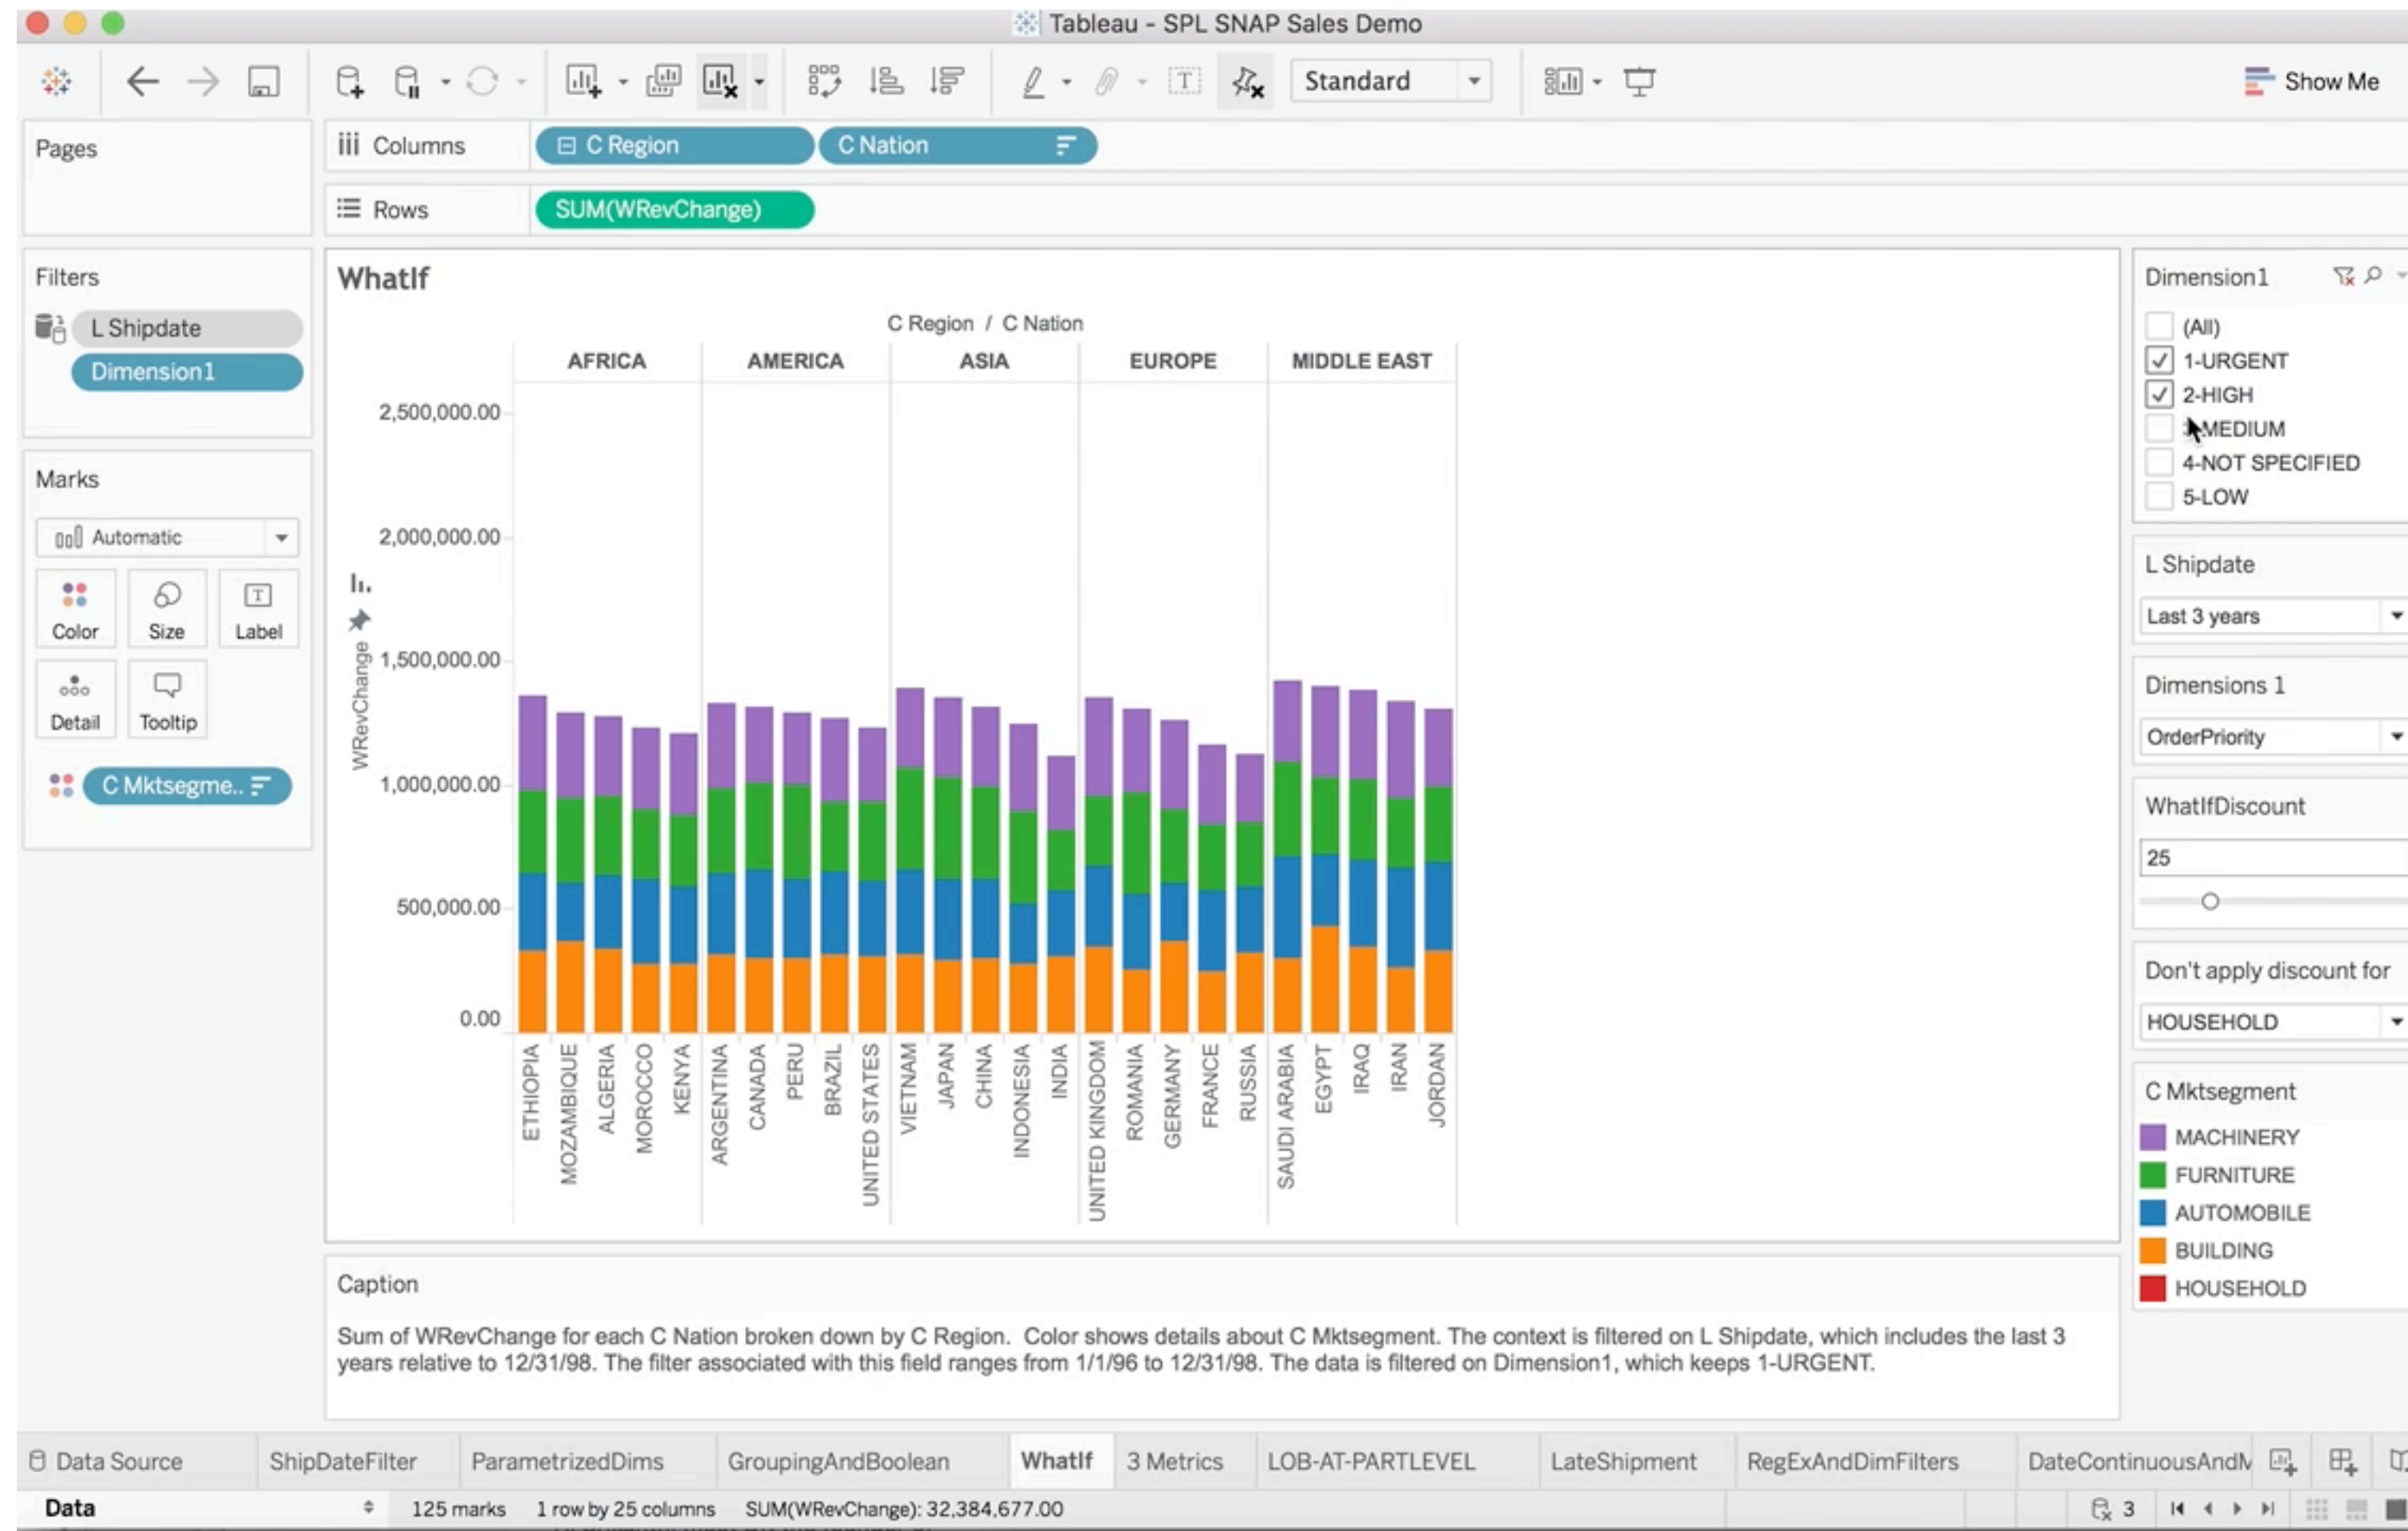
```
SELECT `sales_demo_par`.`c_region` AS  
`calculation_2786672670741028867`,  
AVG(`sales_demo_par`.`l_discount`) AS `avg_l_discount_ok`,  
AVG(`sales_demo_par`.`l_extendedprice`) AS  
`avg_l_extendedprice_ok`, AVG(`sales_demo_par`.`l_quantity`) AS  
`avg_l_quantity_ok`, `sales_demo_par`.`l_linestatus` AS  
`l_linestatus`, SUM(`sales_demo_par`.`l_extendedprice`) AS  
`sum_l_extendedprice_ok`, SUM(`sales_demo_par`.`l_quantity`) AS  
`sum_l_quantity_ok` FROM `demo`.`sales_demo_par`  
`sales_demo_par` WHERE ((CAST(`sales_demo_par`.`l_shipdate`  
AS TIMESTAMP) >= CAST('1996-01-01 00:00:00' AS TIMESTAMP))  
AND (CAST(`sales_demo_par`.`l_shipdate` AS TIMESTAMP) <  
CAST('1999-01-01 00:00:00' AS TIMESTAMP))) GROUP BY  
`sales_demo_par`.`c_region`, `sales_demo_par`.`l_linestatus`
```

Each query in seconds
even when scanning for
millions of rows

demo running on
macbook pro with 8 GB
ram and several apps
running - tableau,
quicktime, keynote etc



What if analysis



adjust discount to simulate how much would revenue have changed if discount was X% less

X is adjustable through the slider

additional filters are added 0 example - for this what-if analysis exclude a given market segment(drop down param)

What IF

```
SELECT  `sales_demo_par`.`c_mktsegment` AS `c_mktsegment`,
        `sales_demo_par`.`c_nation`    AS `c_nation`,
        `sales_demo_par`.`c_region`    AS `c_region`,
        sum(IF((`sales_demo_par`.`c_mktsegment` = 'AUTOMOBILE'),0,((`sales_demo_par`.`l_extendedprice` * (1 -
        (`sales_demo_par`.`l_discount` * 0.63))) - (`sales_demo_par`.`l_extendedprice` * (1 -
        `sales_demo_par`.`l_discount`)))))) AS `sum_revenue_change__copy__ok`
FROM    `demo`.`sales_demo_par` `sales_demo_par`
WHERE    (((cast(`sales_demo_par`.`l_shipdate` AS timestamp) >= cast('1996-01-01 00:00:00' AS timestamp))
AND      (cast(`sales_demo_par`.`l_shipdate` AS timestamp) < cast('1999-01-01 00:00:00' AS timestamp)))
AND      (sales_demo_par`.`c_mktsegment` IN ('AUTOMOBILE', 'FURNITURE'))))
GROUP BY `sales_demo_par`.`c_mktsegment`,
        `sales_demo_par`.`c_nation`,
        `sales_demo_par`.`c_region`
```

spark / hive external tables

sales_demo_par

Query gets rewritten to
use the SNAP index
Up to 100x faster
for slice and dice queries

SNAP OLAP Index

compressed columnar data and index

sales_snap_par

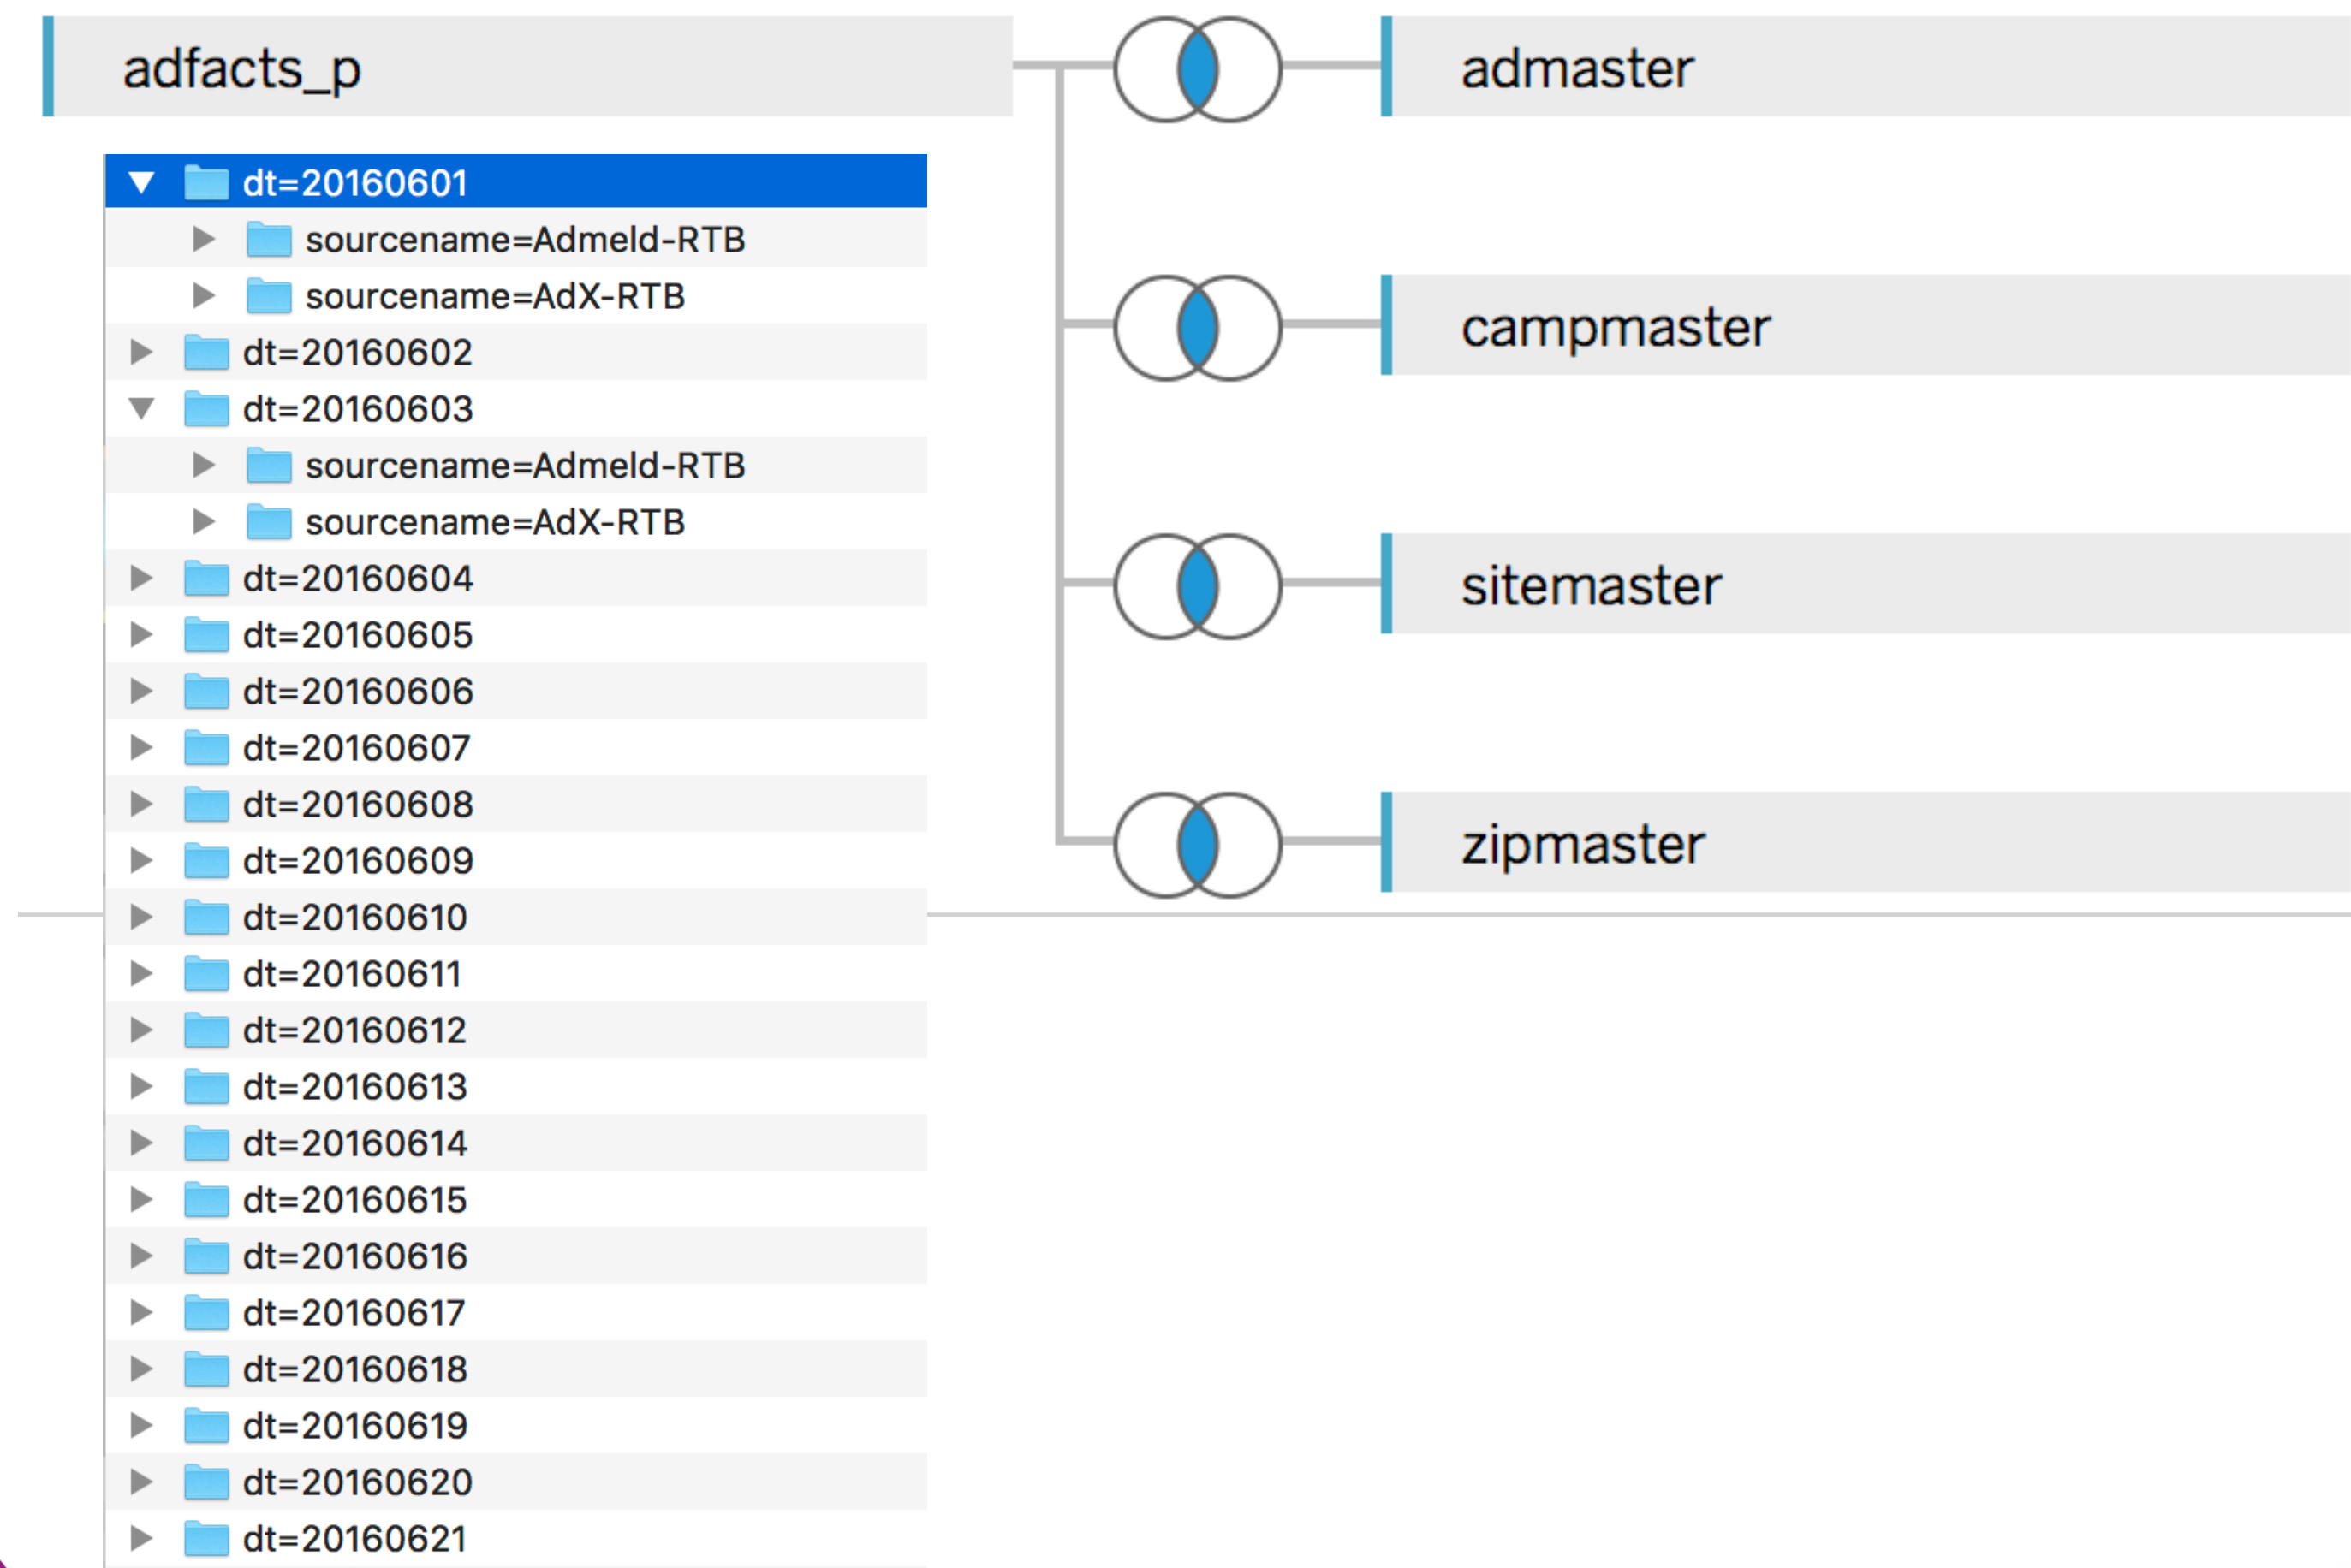
physical data and index in HDFS /S3

```
0 Jan 27 15:33 _SUCCESS
884 Jan 29 08:04 p_year=1992
884 Jan 29 08:10 p_year=1993
884 Jan 29 08:17 p_year=1994
884 Jan 29 08:23 p_year=1995
884 Jan 29 08:29 p_year=1996
884 Jan 29 08:36 p_year=1997
612 Jan 29 08:40 p_year=1998
```



defining star schemas

star schema indexing and querying



1

```
create star schema on adfacts_p as
  many_to_one join of adfacts_p
  with admaster on advertiser_id = advi
  many_to_one join of adfacts_p
  with sitemaster on site_id = siteidmaster
  many_to_one join of adfacts_p
  with campmaster on campaign_id = campid
  many_to_one join of adfacts_p
  with zipmaster on zip = zipcode
```

2

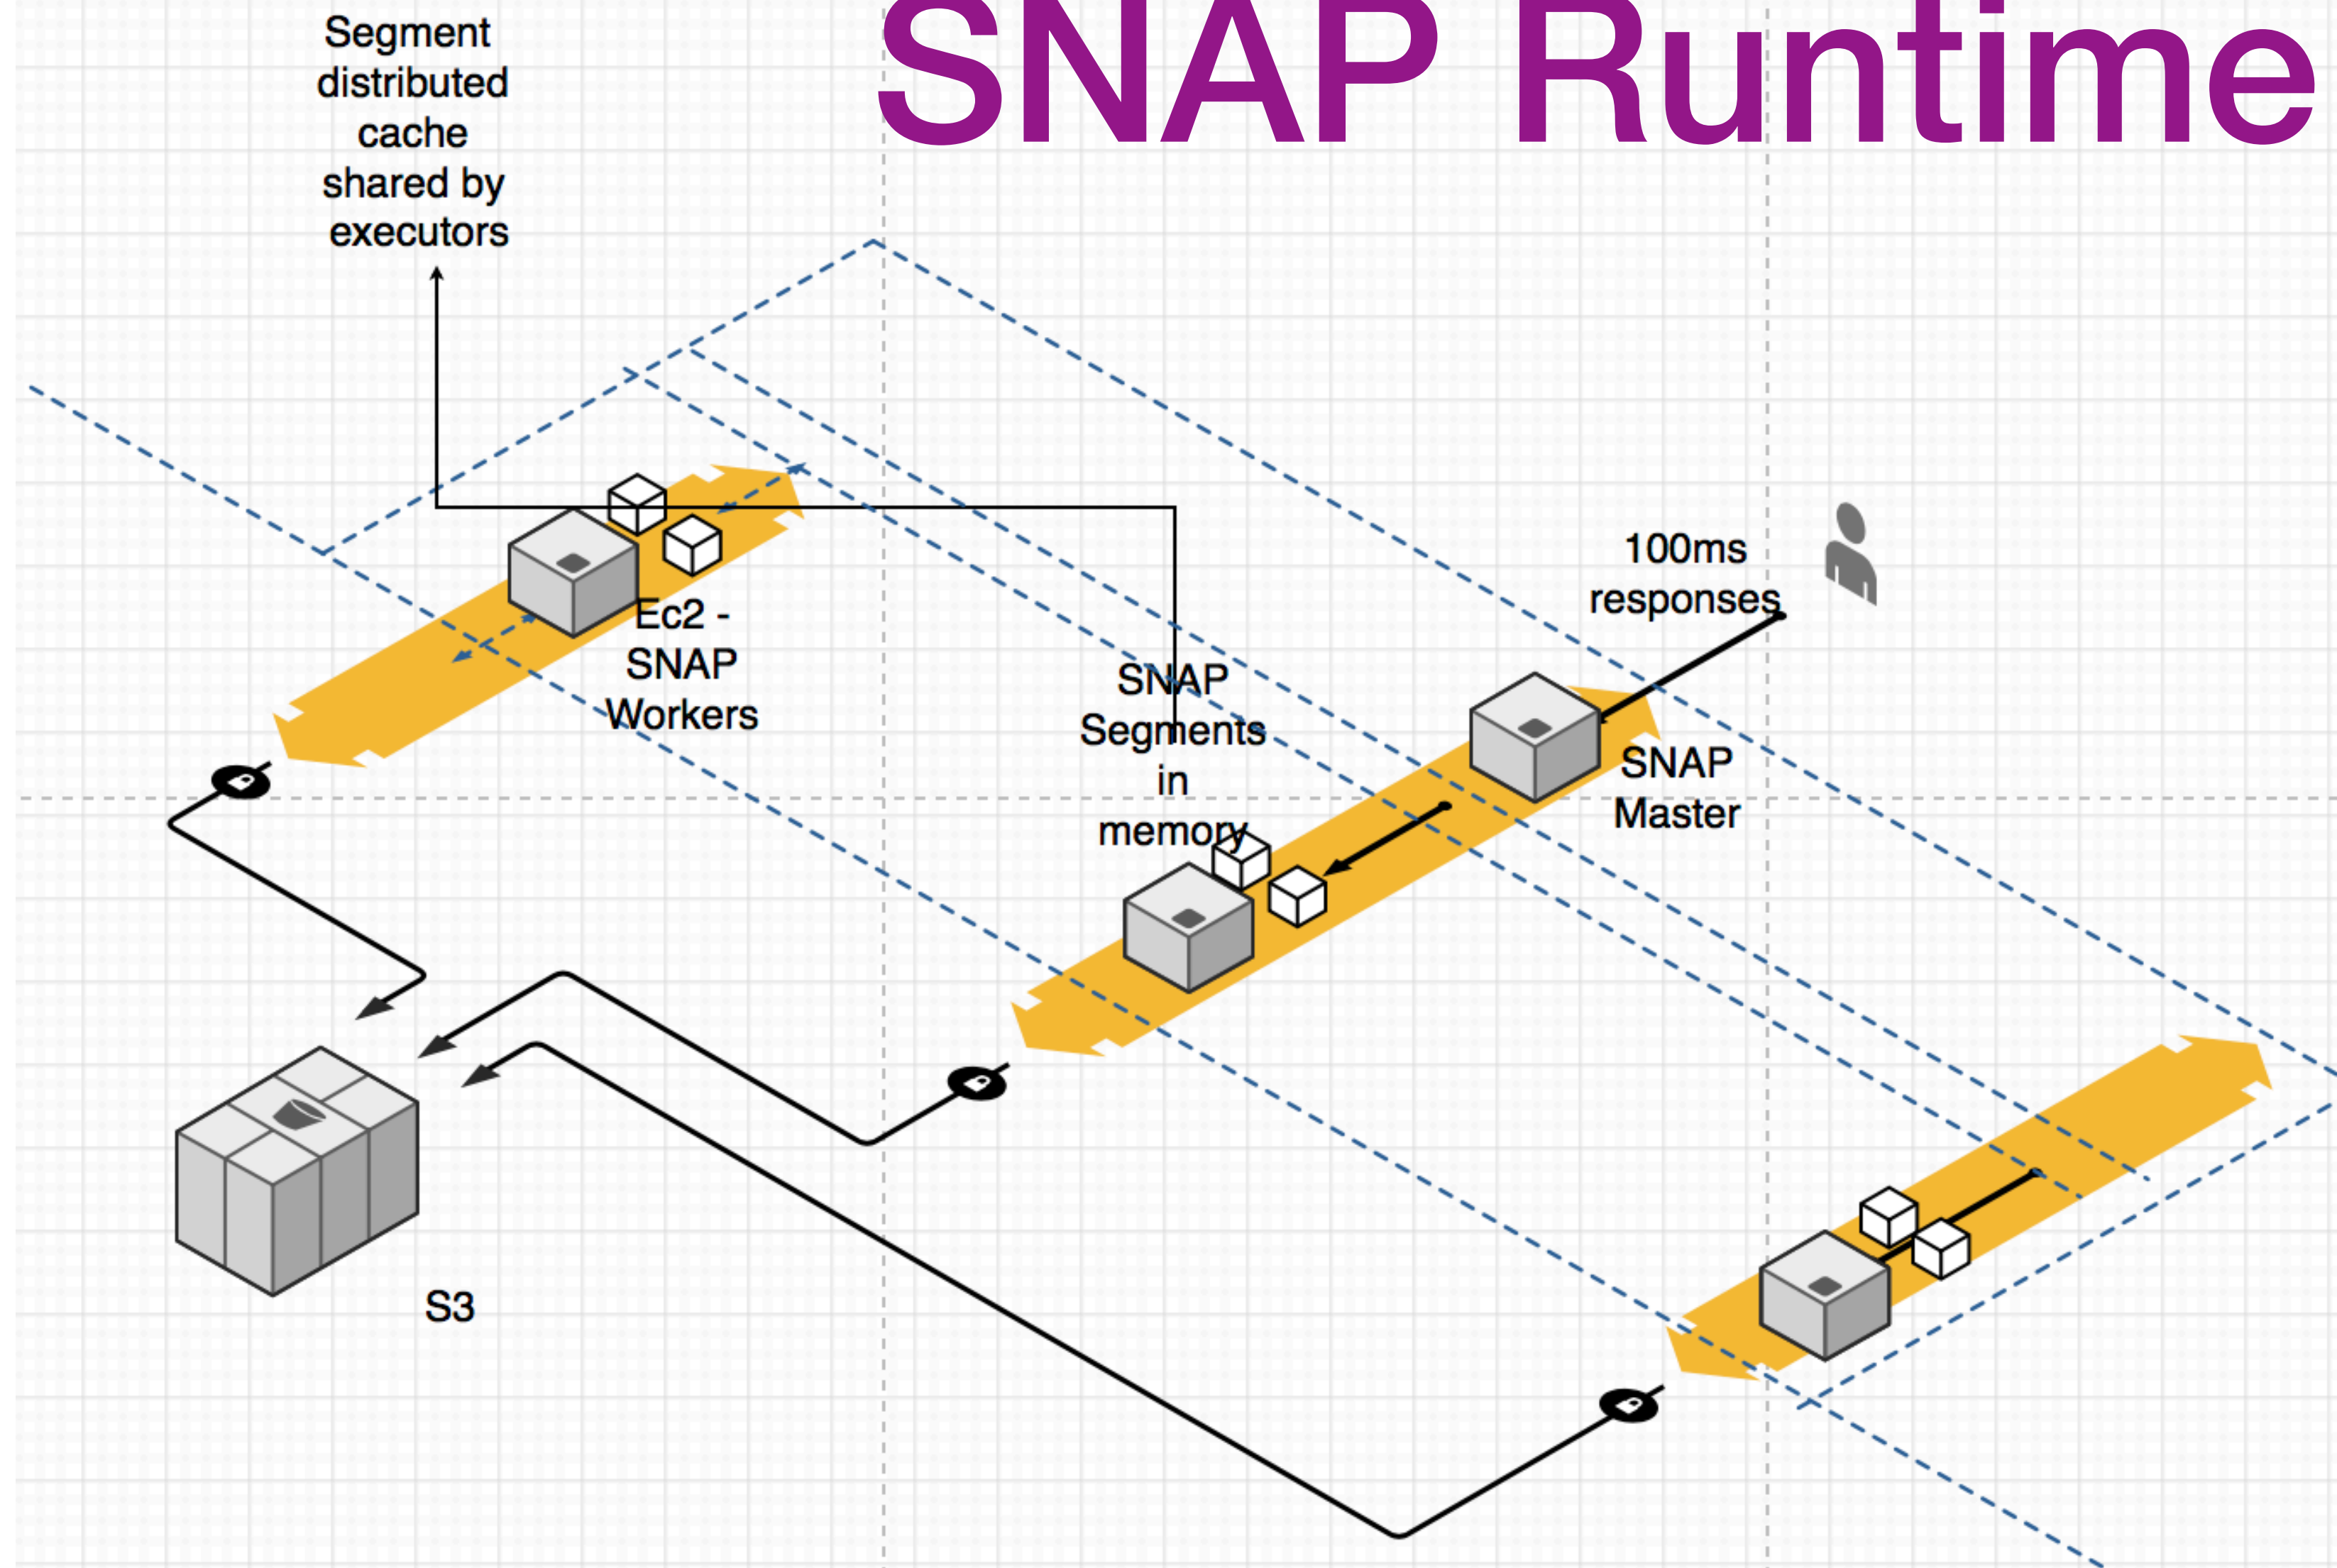
```
create olap index adstarindexp on adfacts_p
  dimension acode is nullable nullvalue "NA"
  dimensions
    "hr,media_type_id,mobile_platform,model,sitenam
    e,advid,advname,campid,campname,zipcode,type,
    primarycity,state,county,timezone,acode,country,e
    stop"
  metrics "revenue,numimpr,numclicks"
  OPTIONS (path "/SNAP/adstarindexP"
    nonAggregateQueryHandling
    "push_project_and_filters",avgSizePerPartition
    "40mb",
    avgNumRowsPerPartition "10000",
    preferredSegmentSize "20mb",
    rowFlushBoundary "1000")
  partition by dt, sourcename;
```

Star Schema

- Once the star schema is defined the process is the same as before for creating an olap index
 - Choose dimensions and metrics across all table columns in the join graph
 - Choose partitioning scheme if needed based on the fact table partitioning
 - Insert into the olap index



SNAP Runtime



- Example - SNAP on S3
- No Hadoop/YARN
- Standalone can support billions of rows and terabytes of data
- SNAP Segments downloaded from s3 to local node distributed cache
- Queries served from segments in-memory
- As new data arrived in S3 SNAP downloads the new segments on request

